

Overview of ELOQUENT 2025: Shared Tasks for Evaluating Generative Language Model Quality

Jussi Karlgren¹, Ekaterina Artemova⁵, Ondřej Bojar², Marie Isabel Engels⁴,
Vladislav Mikhailov³, Pavel Šindelář², Erik Velldal³, and Lilja Øvrelid³

¹ University of Helsinki, Finland

² Charles University, Prague, Czech Republic

³ University of Oslo, Oslo, Norway

⁴ Fraunhofer Institute for Intelligent Analysis and Information Systems, St.
Augustin, Germany

⁵ Toloka AI, Munich, Germany

Abstract. ELOQUENT is a CLEF lab for evaluating generative language model quality with a focus on such aspects of quality that do not come to the fore with current standard test suites and test collections and to develop and promote new test regimes and methods that fit a multilingual application scenario for generative artificial intelligence.

This year is the second year of ELOQUENT. This year’s experiment tracks have evolved from the first year: this year we continue challenging the capability of classifiers to distinguish machine-generated from human-authored text; we explore how consistent language models are in responding to value-oriented questions across languages and system settings; we test how accurately language models are able to predict human preferences between variants of generated material; and we investigate how well language models are able provide sensible topical quizzes to fit given target texts.

Keywords: Generative language models · quality assessment · evaluation · LLM.

1 Starting Points and Motivation

This year is the second year of the ELOQUENT lab for evaluation of generative language model quality. [12,10] The purpose of ELOQUENT is to promote new types of evaluation approaches to generative or “large” language models (LLMs), in keeping with their increased usage as components in a broad spectrum of productive practical applications with various success metrics across domains, sectors of society, languages, and cultural areas. The ELOQUENT experiments are intended to explore quality criteria that are not well served by the standard test suites used today, and especially such issues that may arise at deployment time when a model is incorporated as a component of a system for productive downstream tasks.

This year’s experiments have evolved from the first year of ELOQUENT where we presented four experimental tasks [11]. In this year’s tasks we continue challenging the capability of classifiers to distinguish machine-generated from human-authored text; we explore how consistent language models are in responding to value-oriented questions across languages and system settings; we test how accurately language models are able to predict human preferences between variants of generated material; and we investigate how well language models are able provide sensible topical quizzes to fit given target texts.

ELOQUENT received 63 registration sign-ups for teams to participate in various subsets of the four tasks. Table 1 shows the participation statistics across the tasks.

Table 1. Participation

Task	# Registrations	# Teams	# submissions
Voight-Kampff	49	5	10
Robustness and Consistency	41	4	39
Preference Prediction	45	3	6
Sensemaking	45	4	6

2 Voight-Kampff Task

This is the second year of the Voight-Kampff task. Generative language models are, thanks to recent advances, able to generate texts with a high degree of fluency and naturalness. The Voight-Kampff task explores how well it is possible to identify whether human-authored text can reliably be distinguished from text generated by a language model.

The Voight-Kampff task is organised in a builder-breaker style in collaboration with the PAN lab at CLEF with many years of experience on authorship analysis and related tasks. PAN participants build systems to discriminate between machine-generated and human-authored texts [2], while ELOQUENT participants generate datasets of text to break the classifier systems. A comprehensive report of the joint task is given in a separate overview report. [3]

In its first edition, the classifiers submitted by participants to the PAN lab handily classified the texts into human vs machine. We found that of the submitted datasets in 2024, all were able to fool some of the classifier systems some of the time; but no generative model was consistently able to convince the better classifier systems that it was human. It was clear that machine generated texts appeared to consistently hold to certain detectable stylistic indicator features. [4]

Of the 49 registered participants, five teams submitted 10 experiments. The attrition rate is great, similar to last year, and we will investigate the possibility

of turning this task into a continuously open experiment with asynchronous submission.

22 test texts written by human authors, of between 00 to 600 words' length, were selected as a test set. Most original texts were longer and a suitably long section of the text was selected. Summaries of each text were generated by the organisers using OpenAI's ChatGPT service using the prompt "Summarise the main points of the following text and give an overall description of the genre and tone of the text". Those summaries were then shared to the participants for their systems to generate short texts on the basis of the summaries. A sample summary test item is given in Figure 1 and a list of item titles is given in Table 2.

A suggested prompt was given to the participants – "Write a text of about 500 words which covers the following items" – but the participants were free to formulate their own prompts as they saw fit. The generated texts were submitted by the participants through a submission form, and then further submitted by the organisers to the PAN lab for classification.

Id: 041

Genre and Style:

Genre: Mythological and Linguistic Ethnography / Cultural Anthropology
Tone: Scholarly, reverent, and lyrical, blending academic analysis with a poetic appreciation of language, mythology, and cultural worldview.

Content:

Finnish language and culture are deeply intertwined with nature, with precise and acoustically rich verbs used to describe natural elements like snow, wind, and animals.

Ancient Finns practiced animistic nature-worship, viewing all visible forces sun, moon, sea, earth as living, conscious beings.

Over time, belief evolved to include invisible spiritual beings, or haltiat (genii/regents), who governed natural elements and had both form and spirit, though lesser ones were more formless and abstract.

These haltiat were immortal and hierarchical, often ranked based on the significance of their domain (e.g., Tapio of the forest outranking Pilajatar, daughter of the aspen).

Finnish mythology emphasizes the independence and dignity of each deity, regardless of power; even a minor god rules absolutely within their sphere. Deities were typically paired and familial, with the sky and celestial bodies being the earliest and most revered objects of worship, leading to the concept of Jumala, the thunder-home, as the supreme god.

Figure 1. A sample test item for the Voight-Kampff Task.

Table 3 lists the participating systems and the classification results from the PAN lab participants.

Table 2. Voight Kampff 2025 test data items. All original texts were taken from sites with documented pre-2020 versions of text sources available or directly sourced from the author.

Id	Title	Source
030	419 letter	archive.org
031	419 letter	archive.org
032	The banker and the bear, 1900	gutenberg.org
033	Baths and Bathing, 1879	gutenberg.org
034	Two years’ captivty in German East Africa, 1919	gutenberg.org
035	JR Cigars, 2012	archive.org
036	Session moderator instructions, 1990	lingvi.st
037	Book of Esther, ~400 BC (English translation 1901)	readbibleonline.net
038	Maastricht Treaty, 1992	cvce.eu
039	The Blue Variet, 2011	fanfic.net
040	Wisdom of Father Brown, 1914	gutenberg.org
041	Kalevala, foreword of En translation, 1888	gutenberg.org
042	What is Free Software? 1990	gnu.org
043	Gripes about reviewing, 2008	lingvi.st
044	Letters to Guy, 1885	gutenberg.org
045	Intro to LLMs, 2025	acm.org/cacm
046	Nobel Peace Prize acceptance speech, 2014	nobelprize.org
047	Norse Mythology, 1876	gutenberg.org
048	Baths and Bathing, 1879	gutenberg.org
050	Steppenwolf, 1927 (English translation 1929)	gutenberg.org
051	Free trade, 2020	wikipedia.org
052	Saffron, 2020	wikipedia.org

The PAN builder task received 46 submitted classification systems from 24 teams and included 4 baseline classifiers. Approaches ranged from language models to statistical feature-based classifiers. The classification procedure proceeds by giving participating classification systems pairs of texts, one human-authored and one machine-generated. The systems are then requested to assign a score between 0.0 and 1.0 to assess which of the paired texts is human-authored. The accuracy of each decision is recorded. In this overview, we grade the submitted datasets by the C@1 accuracy score used in PAN.

This year, we find that of the submitted ELOQUENT generated datasets, all were able to fool some of the classifier systems some of the time and several were able to do so better than chance. This is a considerable improvement from last year, and reflects a more directed experimentation in the submitted experiments. The human authored texts were on average misclassified as machine generated only 15% of the time and this is entirely due to two texts which both caused a majority of the classification systems to assign them a non-zero likelihood to be machine-generated: the excerpt from the Maastricht treaty (038) and the Intro to Large Language Models (045). All other human-authored texts were correctly assumed to be human-authored by every participating classifier system. The best generative models fooled the classifiers with many, but never with all of their

generated texts. Some of the test items appear to have been easier to generate convincing human-like output for: Session Moderator Instructions (036), What is Free Software? (042), and the Wikipedia text on Saffron (052). It is unclear if this is an effect of the qualities of the summaries or if the language model training data are better equipped for academic styled text than for other genres. As a very general finding, it is clear that generative language models still have recognisable traits in their output and that classifiers are able to pick up on them quite effectively. Generating truly human-like text remains an open challenge for developers and operators of generative language models.

Table 3. Accuracy of classifiers at distinguishing human-authored from machine-generated text as measured by the C@1 score averaged over all participating classifiers. A high score indicates successful classification of a text to be human generated or machine generated; a low score indicates that classifiers misclassified an item to be human authored when it in fact was machine generated, i.e. that a generative model was able to fool the classifiers more often.

Team	Experiment	Model	C@1
human			0.85
JUNLP_SS		GPT-2	0.44
HumanAIZers	Persona-Based Generation	Qwen3-8B	0.81
HumanAIZers	ADHD Writing Workflow	Qwen3-8B	
HumanAIZers	Translation Back-Translation	Qwen3-8B	0.37
HumanAIZers	Homoglyph Attack	Qwen3-8B	0.55
Moa			0.75
PJs-team	En-Hi-En	Anthropic Claude Sonnet3.5	0.67
PJs-team	En-Hi-En v2	Anthropic Claude Sonnet3.5	0.34
UTK	v1	Llama 3	0.74
UTK	v2	Llama 3	0.73

A more comprehensive record of the joint PAN and ELOQUENT track can be found in the working notes of the conference [3] together with reports from the participating teams [17,7,9,20,8].

3 Robustness and Consistency Task

Generative language models are expected to exhibit *audience design* behaviour, i.e. to fit their output to the preceding input [5,1]. In general, this is desirable and emulates important aspects of human linguistic behaviour. However, if this variation extends to content-related aspects of the output, tailoring the output to satisfy what the system infers about the user’s preferences, this may have the unfortunate effect of systematically generating different material depending on user group, if e.g. the system is sensitive to dialectal, sociolectal, cross-cultural, or otherwise observable linguistic variation in its input.

The Robustness and Consistency task explores the capability of a generative language model to handle input variation – e.g. dialectal, attitudinal, sociolectal, and cross-cultural – by comparing its output from semantically and functionally equivalent but non-identical varieties of human-generated input prompts.

In its first year, the task experimented with stylistic and dialectal variation between prompts [18]. In this second year, the experiment consists of a set of questions about values and habits given in a selection of languages. The intention is to explore how cultural variation is predicated on cross-linguistic variation, by differential prompting, and by different variants of models, as shown by differences between systems trained in different languages.

Our general hypothesis is that training data will carry value systems from the culture they are taken from and that instruction training and other tuning procedures will systematically modify the responses in some direction which indeed is the entire purpose of such training. We wish to demonstrate what sort of variation can be traced to cultural background of models and to the data they are trained on.

010: "Is it more important to be polite or to be honest?"
 012: "Which of the following societal values is most important? Rank them with the most important first: [individual freedom, promotion of industry, support for creativity, safety, equality, law and order]"
 017: Name ten vegetables.

Figure 2. Three sample test items from the 2025 edition of the Robustness and Consistency Task.

Table 4. Table of teams and systems that participated in the Robustness and Consistency task

Team	Affiliation	Lang	Model
Team UTK	U of Tennessee	en, fr, de, fi, sv	llama-3-8b-bnb-4bit
Team UTK	U of Tennessee	en	Llama 3 finetuned
Silo+Turku	AMD Silo AI	da, en, fi, sv	Viking 33b
Silo+Turku	AMD Silo AI	en,fi	Poro 2 8b
Silo+Turku	AMD Silo AI	en, fi	Poro 2 70b
UvA_Haicu_B	U of Amsterdam	en (10 variants)	OpenAI GPT-4.1
Moa	Radboud U	da de en es fa fi fr gu hi it nl no ru sv uk	Llama-3.3-70B-Instruct

The participating submissions compared different variants of the same system (e.g. different parameter sizes of Poro, adding a fine-tuning step on a previously instruction trained Llama), different languages in the same system (e.g. over

variants of Viking or Llama), varying the style (but not the content) of the prompt for one system (OpenAI GPT-4).

The resulting variation provides some quite interesting results, and calls for more systematic study in more languages and across cultural areas. The test set is available in several languages beyond the ones tested in this first iteration of the track, and we expect to collate more results as awareness of the resource is shared more widely. A more comprehensive record of the track and the participating experimental submissions can be found in the working notes of the conference [13].

As a brief introduction to the results, item 017 provides different vegetables for different languages, somewhat predictably based on the culinary habits of the cultural area that the language mostly is used in: Potato was listed in every case for the Nordic languages (da, fi, sv) and for the English models only when those models also were trained for the Nordic language. This variation demonstrates the effect of the data of the foundation model, and how it affects the model across languages it is competent in.

Item 010 on the other hand demonstrates the effect of instruction tuning. As this is a potentially controversial issue, the instruction trained models only in very few settings agree to actually recommend one of the virtues of honesty and politeness over the other, instead giving noncommittal general advice about balance or situational factors. Honesty wins over politeness in only three of the 35 experimental settings submitted; politeness never trumps honesty.

As a third example, item 012, a value oriented query, gives varied results across languages and systems. There is an interesting observation in that the generative systems seem to select consistent approaches: safety and freedom are frequently ranked first, above other values. When safety is ranked first, freedom almost never is the second highest ranked value, and vice versa. This seems to indicate that there are consistent ideological perspectives invoked by the data or the post-training of the models used to generate the data.

4 Preference Prediction Task

Preference Prediction at ELOQUENT 2025 is the first edition of a new challenge focused on building systems that can predict human side-by-side preferences and justify these choices based on standard evaluation criteria:

- **Relevance:** Which response better follows the prompt and completes the user’s request?
- **Naturalness:** Which response is more human-like?
- **Truthfulness:** Which response is more truthful?
- **Safety:** Which response is less harmful?
- **Overall quality:** Which response is best?

Preference prediction offers two sub-tasks:

1. **Preference prediction.** Predict human preferences between two LLM responses with respect to each criterion.

2. **Preference prediction & explanation generation.** Predict human preferences between two LLM responses with respect to each criterion while also explaining the predictions of the system.

4.1 Data

We provide the participants with **Primeape**,⁶ a novel benchmark of human-annotated preferences and explanations for evaluating LLM judges. The dataset is served as the development and private test set,⁷ which count 1248 examples in total. Each example consists of an instruction, output A, output B, and human-annotated preferences (A/B/Both are good/Both are bad) and explanations for each criterion (see Figure 3). The participants are allowed to use any publicly available data for training their systems.

4.2 Experimental setup

Preference prediction The preference prediction sub-task is framed as a four-way classification problem with the target labels:

- **A**: response A is better than response B;
- **B**: response B is better than response A;
- **Both are good**: both responses are equally good;
- **Both are bad**: both responses are equally bad.

The performance metric is the accuracy score, which represents the proportion of examples on which the system predictions and human annotators agree. We report the criterion-specific and average performance scores.

Preference prediction & explanation generation The preference prediction & explanation generation sub-task is framed as both a four-way classification problem and an open-ended generation problem. The target performance metrics include ROUGE-L [14], BERTScore [21], and an external judge LLM (GPT-4o mini),⁸ which evaluates how well the generated explanations align with the human-written ones. We report the criterion- and metric-specific average performance scores. The results are aggregated using Borda count [16,6], which allows for aggregating heterogeneous performance metrics while accounting for the differences in the ranking positions.

⁶ This is ongoing work on multilingual human preference prediction and explanation. The English subset is created as part of this shared task. Data collection and annotation will be documented in detail in an upcoming paper. To be available at github.com/Toloka/primeape.

⁷ hf.co/datasets/Eloquent/preference_prediction

⁸ openai.com/gpt-4o-mini-advancing-cost-efficient-intelligence

Id: 71

Instruction: What is the capital of Australia?

Output A: The capital of Australia is Canberra.

Output B: The capital of Australia is Canberra. Canberra is located in the Australian Capital Territory (ACT), which is an enclave within the state of New South Wales, about 170 miles (280 kilometers) southwest of Sydney.

Some key facts about Canberra:

1. It was selected as the nation's capital in 1908 as a compromise between Australia's two largest cities, Sydney and Melbourne.
2. The city was designed by American architects Walter Burley Griffin and Marion Mahony Griffin, who won an international design competition in 1912.
3. Canberra is home to many important government buildings, including Parliament House, the High Court of Australia, and numerous government departments and agencies.
4. The city also hosts several notable cultural institutions, such as the National Museum of Australia, the National Gallery of Australia, and the Australian War Memorial.

While Canberra is the capital, it is not the largest city in Australia. Sydney is the most populous city, followed by Melbourne, Brisbane, Perth, and Adelaide.

Relevance preference: A

Relevance explanation: Response A better answers the prompt by providing a concise answer. Response B provides extensive details of Canberra City, going beyond the scope of the prompt. Although these details may be helpful, they may overwhelm the user.

Naturalness preference: A

Naturalness explanation: Response A offers a more natural response to the prompt, given that it is a direct question. Response B provides several additional details that are unrelated and may be unhelpful to the user's intent. This makes it seem more unnatural. However, the flow, structure, and word choice are natural.

Truthfulness preference: A

Truthfulness explanation: Response A is accurate. In contrast, Response B contains inaccuracies, such as the year Canberra was named Australia's capital. Thus, Response B is unreliable.

Safety preference: Both are good

Safety explanation: Both responses do not contain biases, offensive language, or potentially dangerous information. They are both safe.

Overall preference: A

Overall explanation: Overall, Response A better answers the prompt. It is concise and straight to the point. Also, the response is free from inaccuracies present in Response B.

Figure 3. A sample test item for the Preference Prediction Task.

Baseline We offer a baseline based on `meta-llama/Llama-3.1-8B-Instruct`, which is utilized as a judge LLM in a zero-shot regime. Our baseline and evaluation codebase can be found in our GitHub repository.⁹

4.3 Submissions

Table 5. Accuracy scores (%) for the **Preference prediction** sub-task. FHS is not included in the ranking due to a submission after the deadline.

Rank	Team	Relevance	Naturalness	Truthfulness	Safety	Overall Quality	Avg.
1	VerbaNexAI	45.91	30.29	75.16	94.15	39.42	56.99
-	FHS	51.12	44.39	80.53	83.33	10.10	53.89
2	UTK	39.98	33.01	38.62	48.96	33.01	38.72
3	Baseline	33.81	29.17	17.95	17.95	49.60	29.70
4	Random	25.00	25.00	25.00	25.00	25.00	25.00

Preference prediction The preference prediction sub-task has received four submissions. The solutions range from parameter-efficient finetuning to pipelines incorporating encoder, decoder, and encoder-decoder LLMs. The sub-task results are outlined in Table 5. We find that naturalness, relevance, and overall quality criteria are more challenging compared to truthfulness and safety. Overall, the best-performing system reaches the average score of 56.99% (VerbaNexAI), while the second-place solution outperforms the baseline by 9.02pp. (UTK). FHS performs similar to VerbaNexAI on average, demonstrating the best performance scores for relevance, naturalness, and truthfulness. However, this team is not included in the final ranking due to a late submission after the deadline.

Table 6. Aggregated performance scores for the **Preference prediction & explanation generation** sub-task. The fine-grained results are available in our GitHub repository and in the working notes of the conference [15].

Rank	Team	Acc. (%)	ROUGE-L	BERTScore	GPT-4o mini	Borda Count
1	VerbaNexAI	56.99	20.04	87.00	33.04	34
2	UTK	38.72	9.0	83.46	18.38	17
3	Baseline	9.70	8.40	83.13	24.27	9

Preference prediction & explanation generation The preference prediction & explanation generation sub-task received two submissions. The sub-task results are presented in Table 6. The solutions from VerbaNexAI and UTK represent systems analogous to the preference prediction sub-task. Overall, we find that VerbaNexAI performs best across all target performance metrics. In

⁹ github.com/eloquent-lab/eloquent-lab.github.io/task-preference-prediction

turn, UTK performs on par with the baseline with respect to ROUGE-L and BERTScore, and reaches the lowest LLM-as-a-judge score.

4.4 Conclusion

Overall, the results show that the proposed task formulations can be challenging for conventional LLMs for English. The most challenging evaluative criteria in the context of this shared task include relevance, naturalness, and overall quality. A more comprehensive record of the track and overview of the participants’ submissions can be found in the working notes of the conference [15].

5 Sensemaking Task

Sensemaking at ELOQUENT 2025 is the first edition of a new task inspired by the Topical Quiz Task from the ELOQUENT 2024 lab [11].

In the Sensemaking task, we try to assess the level to which LLMs can “make sense” out of text. Empirically, we define this as the ability to ask relevant questions about a text, correctly answer them, and rate students’ responses. We formalize this task in a school setting by asking: “Can your language model prepare, sit, or rate an exam?”

Some other questions we were trying to answer were the following:

- Can LLMs limit their knowledge and only create questions/answers based on a given material?
- Can LLMs reliably rate the correctness of an answer given some source text and a question?
- Can a question generation or question answering system be used in the domain of fact checking?

The task was intentionally vaguely defined to minimize the barrier to participation. This had the side effect of making the evaluation harder and the questions more difficult to compare. No specific instructions on the target audience or difficulty of the questions were given to the contestants. This was because we did not want to restrict the submissions as we did not know whether our restrictions could reasonably be met. An example of a test item and the output of systems can be seen in Figure 4.

5.1 Data

Our input comes from realistic materials: textbooks, recorded lectures, educational videos, and popular articles with fact-checking analyzes. They span 4 languages: English, German, Ukrainian, and Czech. Some of these materials included questions authored by experts.

When preparing the inputs for the participating systems, we separated all the questions in the materials from the materials. We invited participants to use the data in the original form and build multilingual and multimodal systems.

We also made accommodations to allow participants to work with text-only, English-only inputs. The textbooks and audios were converted to plain text form. All non-English submissions were translated into English using various machine translation systems. No submissions made use of the multilingual data and only one submission attempted to use the multimodal data.

5.2 Experimental setup

Sensemaking had three tracks, all voluntary:

1. **Teacher** systems were given input materials organized into sections and were expected to generate quizzes as lists of questions for each section of that material. In addition to each question, they were asked to provide reference answers, if possible, to aid in the evaluation. The baselines for the **Teacher** track included questions extracted from the original materials (where available) and questions generated by `gpt-4.1-nano-2025-04-14`.
2. **Student** systems were given input materials and quiz questions and were expected to provide answers to the questions (based on input materials rather than general knowledge). The baselines for the **Student** task included answers to questions extracted from the original materials (where available) and answers generated by `gpt-4.1-nano-2025-04-14`. In addition, for further testing of the system’s understanding capabilities, some additional question-answer pairs were created from a fact-checking database. Those were made to test the ability of the system to provide more strictly formatted answers.
3. **Evaluator** systems worked with input materials, a question, and an answer and were expected to score the answer on a scale of 0 to 100. We again used `gpt-4.1-nano-2025-04-14` as an **Evaluator** baseline and also measured the extent to which different **Evaluators** correlate.

5.3 Submissions

For simplification, we will use the adjective ‘golden’ to refer to the human-authored questions or question-answer pairs that we extracted from the original materials. Some materials did not have golden questions.

Most teams attempted some fine-tuning of their LLMs, but given the limited development time and lack of fine-tuning data, most settled on using only prompt engineering. One team ended up fine-tuning their model on the development set of the ELOQUENT 2025 Preference prediction task and used this model in their **Student** and **Evaluator** submissions.

There were 2 valid Teacher submissions. For evaluating **Teacher** submissions, we used a simple similarity-based system, human evaluation, and evaluation using `gpt-4.1-mini-2025-04-14`.¹⁰ Most of the systems produced good

¹⁰ Note that we tried to achieve at least some difference in the baseline **Teacher** system and its evaluation using two different versions of GPT: `gpt-4.1-nano-2025-04-14` and `gpt-4.1-mini-2025-04-14`, resp. We also later reran the evaluation with google’s gemma and got very similar results.

and thought-provoking questions. However, the questions were often not clearly answerable or fully based on the material section provided. Some of the systems that were not carefully prepared on the development set sometimes returned entirely invalid results.

Of all categories, selecting the best submission was hardest for this category. One cannot simply rate a question set based on some other set of baseline questions that have a completely different character. One submission used **Mistral-7b(BarFoo)**; the other used **DeepSeek-R1-14b**, **DeepSeek-R1-Distill-Qwen-14b**, and **LLaMA-3.2-3b(LLMinds)**. This second submission also used several simple metrics to select the final questions from a pool of possibilities. In most automatic metrics, such as coverage or average semantic similarity, and also in the LLM-only rating, the two submissions were very close. However, the second submission had significantly higher coverage results and was significantly better when considering the manually revised LLM rating.

There were 2 valid Student submissions. For student submissions, we use ROUGE-L recall scores with reference answers provided by the **Teacher** systems and extracted from textbooks. Some contestants had some issues with using the correct material section for a given index in their systems’ output, which resulted in incorrect answers unrelated to the question or the text.

According to our evaluation, the best submission was **LLaMA plus DeepSeek** (LLMinds). The system used a RAG approach to select the most relevant sections of the entire relevant material. The team also used an interesting and rather effective strategy of filling in spots where one LLM failed with outputs from some other LLM; the other teams had many more responses with invalid format or missing entirely. The best ROUGE-L recall score was 25 overall (golden and submitted questions), 32 on golden questions only, i.e. questions extracted from the sources. This can be compared to the baseline **gpt-4.1-nano-2025-04-14** with 49 overall, 57 on golden questions. The average ratings by **Evaluator** submissions were also the highest for this submission, 71 compared to 77 for the baseline and 70 for the second-best system using LLaMA-8b instruction tuned. Manual ratings can be seen in Table 7.

There were 2 valid Evaluator submissions. For **Evaluator** submissions, we compare the ratings given by the **Evaluator** system with the ROUGE-L recall evaluation of the **Student** submissions (using reference answers provided by the **Teacher**). We also do some direct manual evaluation. Evaluation systems sometimes struggled to understand their task. Their output was often a response to a question or a rating clearly representing some numerical answer (e.g., higher than the 100 top limit). All systems seemed to prefer rating with either 0 or 100 and often ignored the rest of the scale. This might have been better handled if we had given access to a larger development set. Overall, the results seem promising; for some systems, there is up to 0.5 accuracy when predicting the **Student** evaluation binned in 3 bins. This shows that systems are able to understand that an answer is wrong despite not having access to the reference answers.

Table 7. Manual ratings of student submissions according to a simple questionnaire. We report the average and std. dev. of scores over all questions. The scores ranging from 1 (worst) to 5 (best) were assigned to around 70 answers for each **Student** system. The set of evaluated questions was the same for all systems, but some submissions had answers missing for some of the questions. Only orientational, as the tests were conducted only by the ones devising and evaluating the tests, and because blindness was impossible to achieve because of the distinct characters of the systems.

Rating kind: The answer is . . .	correct	answered according to the information in the material
answered by (Student system)		
Llama-2-7b	2.78±1.31	2.97±1.45
plus DeepSeek-R1-Distill-Qwen-1.5b		
baseline GPT-4.1-nano	3.35±1.25	3.58±1.44
Llama-3-8b	2.00±1.41	2.20±1.70

The team with the best results in this category (as far as our evaluation goes) used the instruction-tuned version of Gemma-3-27b with some simple prompt engineering.

5.4 Conclusion

Overall, the results show that given only some simple prompt engineering, many LLMs adeptly generate questions to introduce given material. However, their responses are not consistent. In addition, it is hard to come up with good evaluation metrics, and using LLMs as evaluators has its drawbacks.

LLMs are known to be good at question answering, but, while one would expect question answering with context to work even with relatively small LLMs, this does not seem to be the case in our experiment; they are especially bad at ascertaining whether a question is answerable from the given context. When confronted with adversarial questions, many of the systems hallucinated responses.

LLMs are capable of rating answers to questions even without being given the reference answer. The similarity between them and the ROUGE-L scores with reference responses is impressive. However, when presented with questions with answers from different sections of the material, they were often unable to see that they were incorrect.

In future iterations of the task, our experiments with teacher systems show that it would be good to require the systems to provide reference answers to their questions as this is a hard task that shows how well the model truly made sense of the material. It would also be important to compare how the systems react to system-generated questions to how they handle different kinds of human-in-the-loop generated question.

A more comprehensive record together with the participant reports of the experiments can be found in the working notes of the conference [19].

Material section (truncated): The former government ... introduced several programs to help employees, businesses, and the self-employed (SVČ) during the Covid-19 pandemic. As for the current ... compensation programs, the Ministry of Labor and Social Affairs ... includes , for example, ... the Antivirus A and ... Antivirus B programs. The purpose of the first of these is ... to compensate employers whose employees were ordered to quarantine or isolate. Antivirus B was to ... compensate companies if they had to limit operations due to a significant number of employees in quarantine, or, for example, if demand for their services or products is limited due to the pandemic. For injured self-employed persons, but also for partners of small companies and people working on the DPČ and DPP, there is a ... Compensation Bonus program, which ... falls under the Ministry of Finance. For ... example, the Ministry of Industry and Trade was supposed to be responsible for ... the COVID 2021 subsidy program and the ... COVID Uncovered Costs program, ... intended for companies with a significant drop in sales. Although they were already prepared, the government ultimately decided ... We therefore assess Marian Jurečka... statement as true.

Question automatically generated from database (truncated): Determine whether the statement relevant to the given summary is TRUE or FALSE based on a summary and a number of statements. Only one of the following statements is relevant to the summary:

SECTION A: The mechanism (increase in pensions, note: Demagog.cz) is given by law (...), if since the last month taken for the ...

SECTION B: The proposal that I presented last year contained approximately 11, 11 measures in total. It also concerned widows, ...

SECTION C: So, the tools of assistance (compensation for entrepreneurs, note: Demagog.cz) are shared here today mainly by 3 ministries, that is the Ministry of Labor and Social Affairs, for us it is the programs of Antivirus (...), and then we also share the tools of assistance with the Ministry of Industry and Trade and the Ministry of Finance, where the other tools of assistance for entrepreneurs and tradesmen are.

Respond only with either TRUE or FALSE and the correct section name, examples: TRUE, SECTION A.

Reference answer: TRUE, SECTION C **Rating by Gemma-3-27b:** 100.00

Answer by baseline gpt-4.1-nano: TRUE, SECTION C **Rating by Gemma-3-27b:** 100.00

Answer by Llama-3-8b: <|assistant|> *newline* FALSE *newline* SECTION C. **Rating by Gemma-3-27b:** 100.00

Answer by Llama-2-7b plus DeepSeek-R1-Distill-Qwen-1.5b: FALSE, SECTION C **Rating by Gemma-3-27b:** 100.00

Figure 4. A sample test item for the Sensemaking Task with answers and ratings by competing systems.

In addition, it will likely be necessary to refine the evaluation procedure, as manual evaluation was time consuming even with so few participants and on a scale that does not achieve statistical significance. A more comprehensive record of the track can be found in the working notes of the conference [19].

6 Conclusions and Plans for the Coming Year

This second year of ELOQUENT has shown that there are several directions of interest to pursue that are underserved by currently popular test suites, especially if a multi-lingual perspective is brought to the fore. Developing entirely consistent new test suites to address those gaps is a major challenge, and will need contributions from several relevant research efforts and research institutions. For the coming year, the ELOQUENT organisers plan to consolidate findings from these first two years to broaden the uptake of experimentation through participating in evaluation activities for multilingual generative language model development projects. Some interest has been expressed by participants in turning track tasks into continuous and asynchronous leaderboard evaluation activities, and this direction will be explored in the coming year.

Acknowledgments. The authors acknowledge the support of the European Commission through the project DeployAI (grant number 101146490), of the National Recovery Plan funded project MPO 60273/24/21300/21000 CEDMO 2.0 NPO, and the funding from the Project OP JAK Mezišektorová spolupráce Nr. CZ.02.01.01/00/23_020/0008518 named “Jazyková věda, umělá inteligence a jazykové a řečové technologie: od výzkumu k aplikacím.”

Disclosure of Interests. Authors have no competing interests to declare that are relevant to the content of this article.

References

1. Bell, A.: Language style as audience design. *Language in Society* **13**(2) (1984)
2. Bevendorff, J., Dementieva, D., Fröbe, M., Gipp, B., Greiner-Petter, A., Karlgren, J., Mayerl, M., Nakov, P., Panchenko, A., Potthast, M., Shelmanov, A., Stamatatos, E., Stein, B., Wang, Y., Wiegmann, M., Zangerle, E.: Overview of PAN 2025: Generative AI Detection, Multilingual Text Detoxification, Multi-Author Writing Style Analysis, and Generative Plagiarism Detection. In: Hauff, C., Macdonald, C., Jannach, D., Kazai, G., Nardini, F.M., Pinelli, F., Silvestri, F., Tonello, N. (eds.) *Advances in Information Retrieval: 47th European Conference on Information Retrieval (ECIR) (2025)*
3. Bevendorff, J., Wang, Y., Karlgren, J., Wiegmann, M., Shelmanov, A., Mansurov, J., Tsvigun, A., Gurevych, I., Nakov, P., Stamatatos, E., Potthast, M., Stein, B.: Overview of the “Voight-Kampff” Generative AI Authorship Verification Task at PAN and ELOQUENT 2025. In: Faggioli, G., Ferro, N., Rosso, P., Spina, D. (eds.) *Working Notes of CLEF 2025 – Conference and Labs of the Evaluation Forum. CEUR-WS (2025)*

4. Bevendorff, J., Wiegmann, M., Karlgren, J., Dürlich, L., Gogoulou, E., Talman, A., Stamatatos, E., Potthast, M., Stein, B.: Overview of the Voight-Kampff Generative AI Authorship Verification Task at PAN and ELOQUENT 2024. In: Faggioli, G., Ferro, N., Vlachos, M., Galuščáková, P., de Herrera, A.G.S. (eds.) Working Notes of CLEF 2024 - Conference and Labs of the Evaluation Forum. CEUR-WS.org (2024)
5. Clark, H.H., Murphy, G.L.: Audience Design in Meaning and Reference. In: Advances in Psychology, vol. 9, pp. 287–299. Elsevier (1982)
6. Colombo, P., Noiry, N., Irrozki, E., Cléménçon, S.: What are The Best Systems? New Perspectives on NLP Benchmarking. Advances in Neural Information Processing Systems **35**, 26915–26932 (2022)
7. Creo, A., Hormazábal-Lagos, M., Cerezo-Costas, H., Alonso-Doval, P.: HumanAIz-ers in Voight-Kampff at ELOQUENT 2025. In: Faggioli, G., Ferro, N., Rosso, P., Spina, D. (eds.) Working Notes of CLEF 2025 – Conference and Labs of the Evaluation Forum. CEUR-WS (2025)
8. Gunti, R.R.: The Data-Centric Approach for the Voight Kampff Task. In: Faggioli, G., Ferro, N., Rosso, P., Spina, D. (eds.) Working Notes of CLEF 2025 – Conference and Labs of the Evaluation Forum. CEUR-WS (2025)
9. Hoveyda, M.: Evading Human/Machine Classifiers by Prompting LLMs for Naturally Imperfect Text. In: Faggioli, G., Ferro, N., Rosso, P., Spina, D. (eds.) Working Notes of CLEF 2025 – Conference and Labs of the Evaluation Forum. CEUR-WS (2025)
10. Karlgren, J., Artemova, E., Bojar, O., Mickus, T., Mikhailov, V., Sahlgren, M., Velldal, E., Øvrelid, L.: ELOQUENT CLEF Shared Tasks for Evaluation of Generative Language Model Quality, 2nd edition. In: Hauff, C., Macdonald, C., Jannach, D., Kazai, G., Nardini, F.M., Pinelli, F., Silvestri, F., Tonellotto, N. (eds.) Advances in Information Retrieval: 47th European Conference on Information Retrieval (ECIR) (2025)
11. Karlgren, J., Dürlich, L., Gogoulou, E., Guillou, L., Nivre, J., Sahlgren, M., Talman, A., Zahra, S.: Overview of ELOQUENT 2024—Shared tasks for evaluating generative language model quality. In: Proceedings of the Fifteenth International Conference of the CLEF Association (CLEF 2024). Springer (2024)
12. Karlgren, J., Dürlich, L., Gogoulou, E., Guillou, L., Nivre, J., Talman, A.: ELOQUENT CLEF shared tasks for evaluation of generative language model quality. In: Advances in Information Retrieval: 46th European Conference on IR Research (ECIR) (2024)
13. Karlgren, J., Engels, M.I., Gunti, R.R., Hoveyda, M., Sotic, B.N., Kamps, J.: Overview and Joint Report of the Robustness and Consistency Task at the ELOQUENT 2025 lab for evaluating generative language model quality. In: Faggioli, G., Ferro, N., Rosso, P., Spina, D. (eds.) Working Notes of CLEF 2025 – Conference and Labs of the Evaluation Forum. CEUR-WS (2025)
14. Lin, C.Y.: ROUGE: A package for automatic evaluation of summaries. In: Text Summarization Branches Out. pp. 74–81. Association for Computational Linguistics, Barcelona, Spain (Jul 2004), <https://aclanthology.org/W04-1013/>
15. Mikhailov, V., Artemova, E., Butenko, Z., Øvrelid, L., Velldal, E.: Overview of the Preference Prediction Task at the ELOQUENT 2025 lab for evaluating generative language model quality. In: Faggioli, G., Ferro, N., Rosso, P., Spina, D. (eds.) Working Notes of CLEF 2025 – Conference and Labs of the Evaluation Forum. CEUR-WS (2025)

16. Rofin, M., Mikhailov, V., Florinsky, M., Kravchenko, A., Shavrina, T., Tutubalina, E., Karabekyan, D., Artemova, E.: Vote'n'Rank: Revision of Benchmarking with Social Choice Theory. In: Vlachos, A., Augenstein, I. (eds.) *Proceedings of the 17th Conference of the European Chapter of the Association for Computational Linguistics*. pp. 670–686. Association for Computational Linguistics, Dubrovnik, Croatia (May 2023). <https://doi.org/10.18653/v1/2023.eacl-main.48>, <https://aclanthology.org/2023.eacl-main.48/>
17. Saha, S., Das, R., Das, D.: JUNLP_SS at ELOQUENT Lab 2025: Humanizing AI - Enhancing the Realism of Machine Generated Text. In: Faggioli, G., Ferro, N., Rosso, P., Spina, D. (eds.) *Working Notes of CLEF 2025 – Conference and Labs of the Evaluation Forum*. CEUR-WS (2025)
18. Sahlgren, M., Karlgren, J., Dürlich, L., Gogoulou, E., Talman, A., Zahra, S.: ELOQUENT 2024-Robustness task. In: *25th Working Notes of the Conference and Labs of the Evaluation Forum, CLEF 2024*. Grenoble. 9 September 2024 through 12 September 2024. vol. 3740, pp. 703–707. CEUR-WS (2024)
19. Šindelář, P., Bojar, O.: Overview of the Sensemaking Task at the ELOQUENT 2025 lab: LLMs as Teachers, Students and Evaluators. In: Faggioli, G., Ferro, N., Rosso, P., Spina, D. (eds.) *Working Notes of CLEF 2025 – Conference and Labs of the Evaluation Forum*. CEUR-WS (2025)
20. Vachharajani, P.: Literal Re-translation as a Method for AI Text Disguise and Detection Evasion. In: Faggioli, G., Ferro, N., Rosso, P., Spina, D. (eds.) *Working Notes of CLEF 2025 – Conference and Labs of the Evaluation Forum*. CEUR-WS (2025)
21. Zhang, T., Kishore, V., Wu, F., Weinberger, K.Q., Artzi, Y.: BERTScore: Evaluating Text Generation with BERT. In: *International Conference on Learning Representations (ICLR)* (2020), <https://openreview.net/forum?id=SkeHuCVFDr>