

ELOQUENT

EVALUATION OF GENERATIVE LANGUAGE MODEL QUALITY AND
CHARACTERISTICS

Task 1: Voight-Kampff

Jussi Karlgren
University of Helsinki & AMD Silo AI

Task 1: Voight-Kampff – Has a machine or a human written this?

Can we differentiate human-generated from AI-generated text?

- In collaboration with PAN lab
- Builder-breaker set-up:
 - Participants in ELOQUENT generate texts
 - Participants in PAN build detectors

<https://eloquent-lab.github.io/>



Can your LLM fool a classifier
into thinking it is human?

<https://eloquent-lab.github.io/>

Example generation topic for Voight-Kampff task

Content:

- Finnish language and culture are deeply intertwined with nature, with precise and acoustically rich verbs used to describe natural elements like snow, wind, and animals.
- Ancient Finns practiced animistic nature-worship, viewing all visible forces sun, moon, sea, earth as living, conscious beings.
- Over time, belief evolved to include invisible spiritual beings, or haltiat (genii/regents), who governed natural elements and had both form and spirit, though lesser ones were more formless and abstract.
- These haltiat were immortal and hierarchical, often ranked based on the significance of their domain (e.g., Tapio of the forest outranking Pihlajatar, daughter of the aspen). Finnish mythology emphasizes the independence and dignity of each deity, regardless of power; even a minor god rules absolutely within their sphere.
- Deities were typically paired and familial, with the sky and celestial bodies being the earliest and most revered objects of worship, leading to the concept of Jumala, the thunder-home, as the supreme god.

Genre:

- Mythological and Linguistic Ethnography / Cultural Anthropology

Tone:

- Scholarly, reverent, and lyrical, blending academic analysis with a poetic appreciation of language, mythology, and cultural worldview.

Voight-Kampff submissions

Team	Experiment	Model	C@1 ⁻¹	FPR	FNR	Ref.
<i>Human</i>			0.136	0.148	–	
JUNLP_SS		GPT-2	0.477	–	0.465	[10]
HumanAlzers	Translation Back-Translation	Qwen3-8B	0.618	–	0.609	[11]
HumanAlzers	Homoglyph Attack	Qwen3-8B	0.436	–	0.428	
HumanAlzers	ADHD Writing Workflow	Qwen3-8B	0.142	–	0.136	
HumanAlzers	Persona-Based Generation	Qwen3-8B	0.083	–	0.081	
Moa			0.178	–	0.175	[14]
PJS-team	En-Hi-En v2	Claude Sonnet3.5	0.626	–	0.625	[12]
Pjs-team	En-Hi-En	Claude Sonnet3.5	0.280	–	0.276	
UTK	v2		0.222	–	0.214	[13]
UTK	v1		0.177	–	0.172	

Observations

- Better results this year! ("better" = generative systems fool classifiers)

Observations

- Better results this year! ("better" = generative systems fool classifiers)
- All experimental submissions fooled some of the classifiers some of the time
- Several were able to do so better than chance
- None fooled classifiers with all generated texts
- Approaches include translation, norm departure generation in postprocessing, and creative prompting

Observations

- Better results this year! ("better" = generative systems fool classifiers)
- All experimental submissions were able to fool some of the classifier systems some of the time
- Several were able to do so better than chance
- None were able to fool classifiers with all generated texts
- Approaches include translation, norm departure generation in postprocessing, and creative prompting
- **Human authored texts were on average misclassified as machine generated only 15% of the time**
- Entirely due to two texts
 - excerpt from the Maastricht treaty (038)
 - intro to Large Language Models (045)
- Every other human-authored text was correctly assumed to be human by every classifier

Observations

- Better results this year! ("better" = generative systems fool classifiers)
- All experimental submissions were able to fool some of the classifier systems some of the time
- Several were able to do so better than chance
- None was able to fool classifiers with all generated texts
- Approaches include translation, norm departure generation in postprocessing, and creative prompting
- Human authored texts were on average misclassified as machine generated only 15% of the time
- Entirely due to two texts
 - excerpt from the Maastricht treaty (038)
 - intro to Large Language Models (045)
- Every other human-authored text was correctly assumed to be human by every classifier
- **Some items seemed to be easier to emulate human authorship for**
 - Conference Session Moderator Instructions (036)
 - What is Free Software? (042)
 - Wikipedia text on Saffron (052)

Observations

- Better results this year! ("better" = generative systems fool classifiers)
- All experimental submissions were able to fool some of the classifier systems some of the time
- Several were able to do so better than chance
- None was able to fool classifiers with all generated texts
- Approaches include translation, norm departure generation in postprocessing, and creative prompting
- Human authored texts were on average misclassified as machine generated only 15% of the time
- Entirely due to two texts
 - excerpt from the Maastricht treaty (038)
 - intro to Large Language Models (045)
- Every other human-authored text was correctly assumed to be human by every classifier
- Some items seemed to be easier to emulate human authorship for
 - Conference Session Moderator Instructions (036)
 - What is Free Software? (042)
 - Wikipedia text on Saffron (052)
- **Generating truly human-like text remains a challenge!**

No.

Can your LLM fool a classifier
into thinking it is human?

<https://eloquent-lab.github.io/>

Attrition rate better than last year but still not great

49 registrations: 5 teams submitted 10 experiments.

How can we get more submissions?

The task is quite simple to to run, and allows for some fun experimentation!

Is there some tweak we can do to make it easier?

Could the task move to a continuous leaderboard?

Is there a validation partner we could engage with?